Global Academy of Technology

Autonomous Institute, Affiliated to VTU

GLOBAL ACADEMY OF TECHNOLOGY
growing ahead of time
Autonomous Institute, Affiliated to VTU

Department of Artificial Intelligence & Machine Learning

Ability Enhancement Course (22AML67A)

On

"Advanced Pneumonia Classification System"

Submittted By

Team member-1
USN: 1GA22AI032
NAME: Pavan kumar R

Team member-2
USN: 1GA22AI053
NAME: Suraj.G

Team member-3
USN: 1GA23AI404
NAME: Shrinivas

Under the Guidance of
Prof. Lakshmidevi H M
Assistant Professor
Department of AI & ML

PROBLEM STATEMENT

Pneumonia is a serious respiratory infection that affects millions of people around the world. It remains one of the leading causes of illness and death, especially among vulnerable groups like infants, the elderly, and those with weakened immune systems. Global health reports show that pneumonia causes more deaths in children under five than any other infectious disease. The impact of this disease is even greater in low- and middle-income countries, where access to timely medical care is often limited.

Early detection and correct diagnosis of pneumonia are vital to reducing complications, preventing deaths, and ensuring proper treatment. Traditionally, chest X-rays are a main tool that radiologists and clinicians use to identify pneumonia's presence and severity. However, interpreting these X-rays manually comes with several challenges. The diagnostic process can take a long time, especially in stressed healthcare systems, and is prone to human error due to fatigue, subjective judgment, or limited experience. Additionally, the visual signs of pneumonia on chest X-rays—like opacities, consolidation, and infiltrates—can sometimes look similar to other conditions, such as tuberculosis, lung cancer, or COVID-19, making it harder to differentiate even for skilled professionals.

Another major issue is the inconsistent availability of expert radiologists, especially in remote or underdeveloped areas. This can lead to delays in diagnosis and treatment, raising the risk of severe health problems. Furthermore, image quality can vary a lot due to differences in equipment, patient movement, or poor imaging conditions, which complicates the diagnostic process even more.

Given these challenges, there is a strong need for a reliable, efficient, and scalable solution to help healthcare professionals diagnose pneumonia more accurately and quickly. An AI- powered system that uses deep learning techniques to analyze chest X-ray images shows great potential to fill these gaps. This system can quickly process large numbers of images, perform consistently, and act as a helpful diagnostic tool in both clinical and remote settings. Additionally, when combined with machine learning classifiers and smart medical assistants, these tools can not only detect pneumonia but also provide probability scores, visualize findings, and explain predictions to aid clinical decision-making.

Developing an automated pneumonia classification system is a significant step toward modernizing healthcare diagnostics, improving patient outcomes, and making expert-level medical insights available in various geographic and economic contexts.

Table of Content

INTRODUCTION

The Advanced Pneumonia Classification System is a smart diagnostic platform that addresses the key challenges in pneumonia detection using chest X-ray imaging. Traditional diagnostic methods rely heavily on expert radiologists and can take a lot of time while also being prone to mistakes, especially in areas with limited healthcare resources. This system offers a hybrid AI-powered solution that combines deep learning and classical machine learning algorithms to automate and improve the diagnostic process.

At the heart of the system is a Convolutional Neural Network (CNN), specifically the EfficientNetB0 model, which has been trained using transfer learning. This model extracts deep features from chest X-ray images. These features capture complex patterns in medical imaging. They pass through dense neural layers with dropout regularization before going to several strong machine learning classifiers. These classifiers include Gradient Boosting Machines (GBM), Random Forest (RF), and Support Vector Machines (SVM). They make the final diagnosis by analyzing the learned features and providing probability-based predictions of pneumonia presence.

What makes this system special is its smart integration with llm models, a modern language model that serves as a medical assistant. llm models offers real-time, human-like interactions for healthcare workers and patients by answering questions about pneumonia, explaining diagnostic results, and generating professional-style reports with risk assessments. This feature adds extra usability and clinical relevance, helping bridge the gap between technical results and medical interpretation.

The front end is created using Streamlit, which provides a clean, interactive interface. Users can upload images, set training parameters, visualize performance metrics like ROC curves and confusion matrices, and interact with the chatbot assistant. The integration of TensorFlow and Scikit-learn on the backend ensures high performance in model training and smooth deployment of traditional machine learning algorithms. Together, this system improves the speed and accuracy of pneumonia detection and promotes accessibility, transparency, and support in clinical decision-making, making it an excellent tool for medical institutions and telehealth settings.

SYSTEM REQUIREMENTS

The Advanced Pneumonia Classification System is designed to run efficiently on various computing devices, making it accessible for both developers and healthcare professionals. While deep learning applications typically require high-end hardware, this system is optimized for performance without sacrificing accuracy or user experience.

The minimum recommended RAM is 4GB, which is enough to run the trained model and conduct inference tasks, such as analyzing chest X-ray images and generating results. This allows the system to work well even on mid-range laptops and desktops. However, for users looking to train or retrain the deep learning model with new or custom datasets, it's best to use a system with a dedicated GPU (Graphics Processing Unit). GPUs speed up training by parallelizing matrix operations central to deep learning algorithms. Without a GPU, training on large datasets can take several hours or even days.

The system is built using Python, and its core functionality relies on several open-source libraries and frameworks:

TensorFlow and Keras are used for building, training, and managing the deep learning models, specifically for implementing and fine-tuning the EfficientNetB0 backbone. These libraries offer many tools for deep learning processes, such as automatic differentiation, model serialization, and GPU acceleration.

Scikit-learn is used for creating and evaluating traditional machine learning classifiers like Random Forest, SVM (Support Vector Machine), and Gradient Boosting Machine. It also helps with data preprocessing, model evaluation, and performance analysis.

Streamlit is a lightweight and interactive web application framework that allows the entire system to be launched as a web app. This feature enables real-time image uploads, displays diagnosis results, provides interactive charts (such as ROC curves and confusion matrices), and even facilitates chatbot communication, all from a browser interface with no frontend coding needed.

LLM models is incorporated to offer advanced medical language processing features. This includes generating explanations about pneumonia, responding to user queries in natural language, and aiding in automated report generation with clinical insights.

IMPLEMENTATION

The Pneumonia Classification System uses a clear and efficient AI process designed to classify chest X-ray images accurately and understandably. The process starts with image preprocessing, which is essential for preparing raw X-ray data for the model. This phase includes several steps:

Resizing images to a uniform size of 224x224 pixels to meet the input size needs of deep learning models.

Normalization, where pixel intensity values are adjusted to a standard range (usually between 0 and 1) to help the model train well and converge faster.

Data augmentation is used to increase dataset variety and prevent overfitting. Techniques like random rotation, zooming, flipping, and contrast adjustment mimic real-world differences in X-rays, which helps the model generalize better to new data.

After preprocessing, the images go into a deep learning feature extractor. Specifically, EfficientNetB0, which is a lightweight but effective convolutional neural network (CNN) architecture that has demonstrated top performance in various medical imaging tasks.
EfficientNetB0 uses compound scaling to balance depth, width, and resolution. This approach enables it to extract detailed, layered features from images with fewer parameters and lower computational costs.

The features extracted are then processed through several fully connected dense layers, where further transformations occur. These layers use dropout regularization to avoid overfitting by randomly turning off some neurons during training. Once these dense representations are ready, they go into traditional machine learning classifiers like Gradient Boosting Machine (GBM), Random Forest (RF), and Support Vector Machine (SVM). These classifiers are known for their reliability, interpretability, and performance on structured data, and they help make the final binary decision about whether pneumonia is present.

The user interface of the system is created using Streamlit, a fast and interactive Python web framework suited for machine learning applications. The web app enables users to upload chest X-ray images in common formats such as JPG, JPEG, or PNG. After uploading an image, the system processes it through the model and shows classification probabilities (for example, 92% chance of pneumonia).

Additionally, the interface offers real-time training visualizations that include loss and accuracy graphs, ROC curves, confusion matrices, and probability distributions. These tools assist developers, clinicians, and data scientists in evaluating and understanding the model's behavior and performance interactively.

A notable feature is the LLM-powered chatbot. This AI assistant allows users to ask medical questions about pneumonia and receive intelligent, medically relevant answers. It can also create automated reports summarizing diagnostic results, giving users a clear overview of outcomes and risk factors.

In summary, this implementation combines deep learning for image analysis, traditional machine learning for decision-making, and an interactive frontend for usability. This results in a very effective and deployable AI solution for detecting pneumonia.

SCREENSHOTS

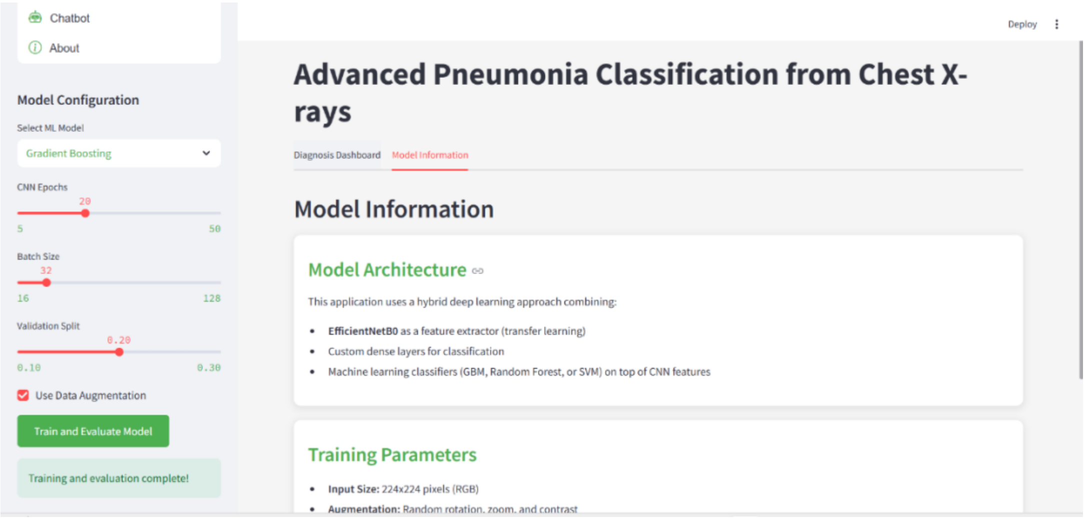

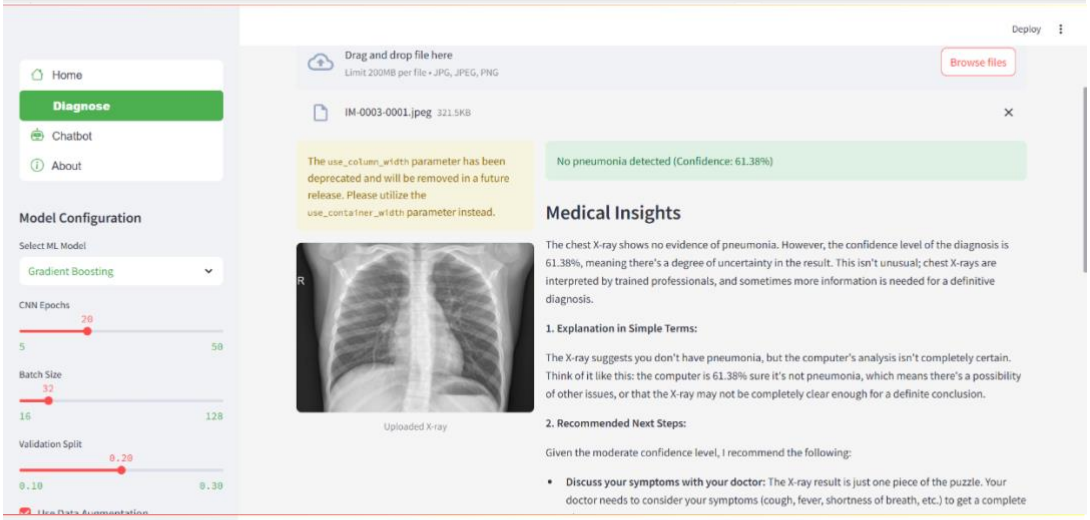

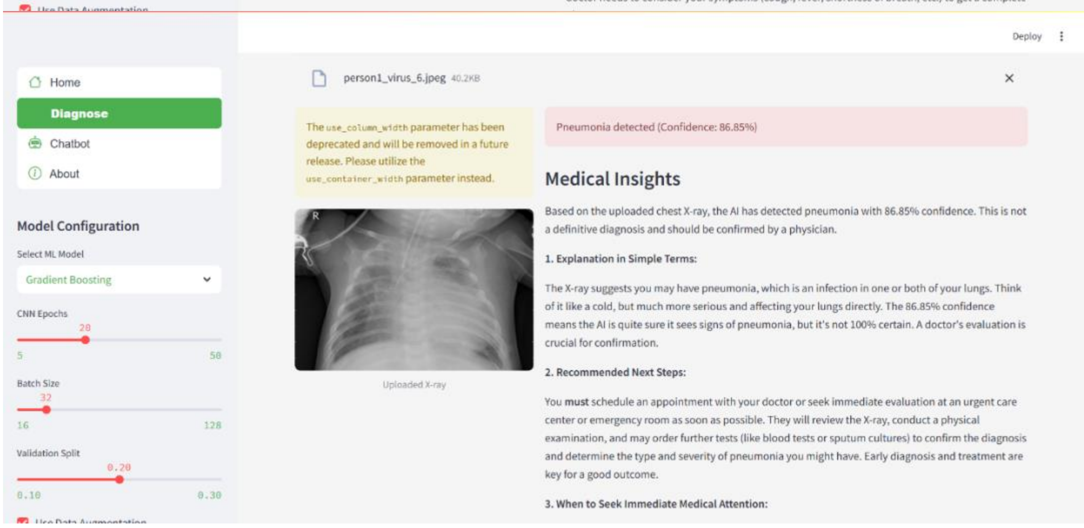

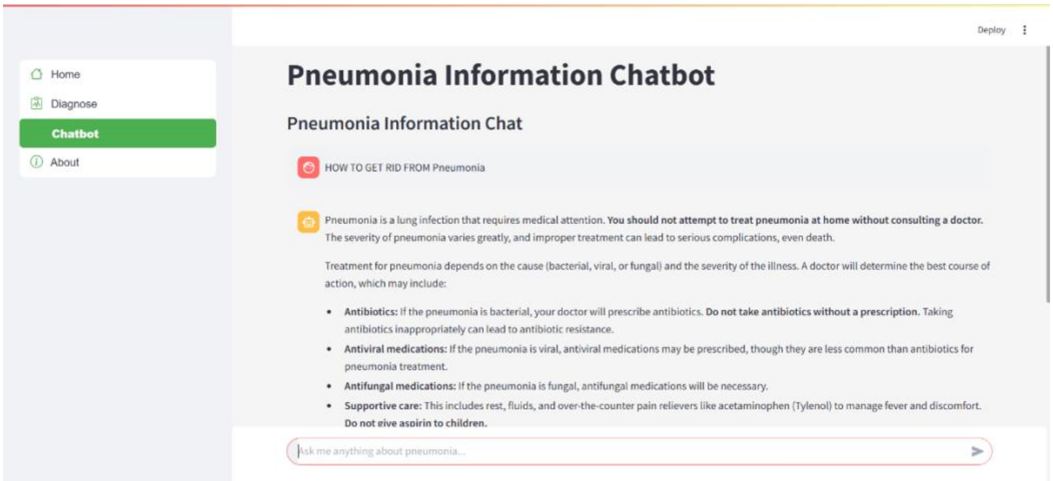
Deploy
Home
Diagnose
Chatbot
About
Pneumonia Information Chatbot
Pneumonia Information Chat
HOW TO GET RID FROM Pneumonia
Pneumonia is a lung infection that requires medical attention. You should not attempt to treat pneumonia at home without consulting a doctor. The severity of pneumonia varies greatly, and improper treatment can lead to serious complications, even death.
Treatment for pneumonia depends on the cause (bacterial, viral, or fungal) and the severity of the illness. A doctor will determine the best course of action, which may include:
Antibiotics: If the pneumonia is bacterial, your doctor will prescribe antibiotics. Do not take antibiotics without a prescription. Taking antibiotics inappropriately can lead to antibiotic resistance.
Antiviral medications: If the pneumonia is viral, antiviral medications may be prescribed, though they are less common than antibiotics for pneumonia treatment.
Antifungal medications: If the pneumonia is fungal, antifungal medications will be necessary.
Supportive care: This includes rest, fluids, and over-the-counter pain relievers like acetaminophen (Tylenol) to manage fever and discomfort. Do not give aspirin to children.
Ask me anything about pneumonia...

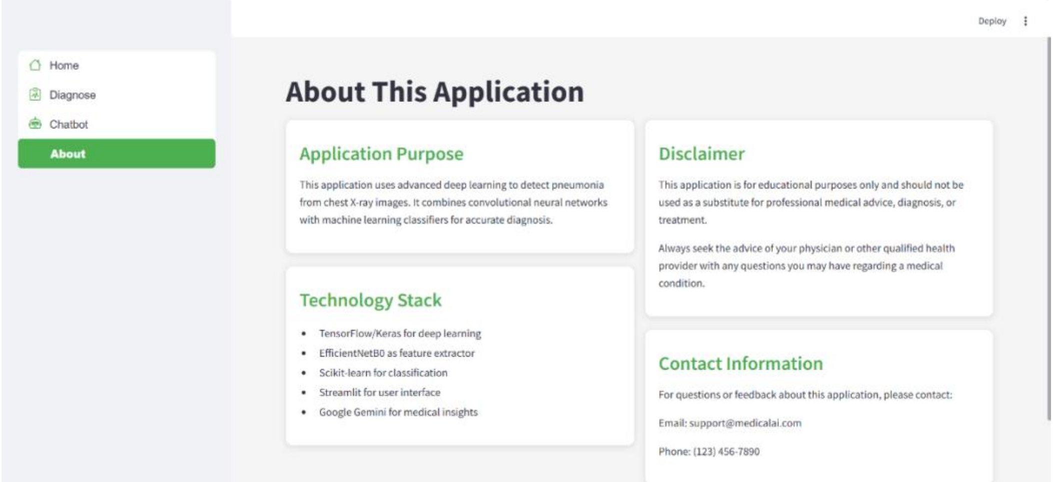
Deploy
Home
Diagnose
Chatbot
About
About This Application
Application Purpose
This application uses advanced deep learning to detect pneumonia from chest X-ray images. It combines convolutional neural networks with machine learning classifiers for accurate diagnosis.
Disclaimer
This application is for educational purposes only and should not be used as a substitute for professional medical advice, diagnosis, or treatment.
Always seek the advice of your physician or other qualified health provider with any questions you may have regarding a medical condition.
Technology Stack
TensorFlow/Keras for deep learning
EfficientNetB0 as feature extractor
Scikit-learn for classification
Streamlit for user interface
Google Gemini for medical insights
Contact Information
For questions or feedback about this application, please contact:
Email: support@medicalai.com
Phone: (123) 456-7890

CONCLUSION AND FUTURE SCOPE

The Advanced Pneumonia Classification System is a reliable, smart, and easy-to-use diagnostic aid. It aims to help in the early detection and management of pneumonia. Its hybrid design combines deep learning models like EfficientNetB0 with traditional machine learning classifiers such as Random Forest, GBM, and SVM. This combination provides high accuracy while remaining interpretable. The system can automatically process and analyze chest X-rays, making it a valuable second-opinion tool for doctors, especially in areas where access to experienced radiologists is limited. In telemedicine, healthcare professionals often need to make quick decisions from a distance. This AI tool can flag potential pneumonia cases with confidence scores, allowing for quicker intervention and treatment.

Additionally, the system works with LLM model to offer interactive explanations. It provides not only predictions but also medical insights in plain language. This helps healthcare providers understand the reasoning behind the AI's predictions and builds trust in the system. In rural and underserved regions, where diagnostic resources are often lacking, this tool can run on low-cost hardware or be integrated with mobile platforms to extend diagnostic services beyond urban hospitals.

Looking ahead, the system is built for growth and change. One major improvement being considered is adding the ability to classify other chest diseases like tuberculosis, lung cancer, and COVID-19 pneumonia, which can have similar radiographic features. This would change the tool from focusing on one disease to becoming a complete thoracic imaging assistant, greatly increasing its clinical usefulness.

Moreover, adding DICOM (Digital Imaging and Communications in Medicine) support will allow the system to connect easily with hospital Picture Archiving and Communication Systems (PACS). This step is vital for clinical use because it enables the system to integrate into existing hospital workflows. Radiologists will then have access to AI insights without leaving their main diagnostic platforms.

Another exciting possibility is developing a mobile application that lets clinicians or trained community health workers capture and upload X-rays using portable imaging devices. When combined with cloud-based processing and LLM-powered support, the mobile platform could provide AI-driven diagnostics in the field, speeding up responses during health emergencies or outreach missions.

In all future developments, the system will focus on ethical aspects, including patient data privacy, reducing bias in training data (to ensure fairness across different groups), and meeting regulatory standards like HIPAA and GDPR. By balancing technological progress with ethical responsibilities, the Advanced Pneumonia Classification System aims to become a trusted, clinically relevant tool that empowers healthcare providers and enhances patient outcomes worldwide.